

Choosers: The design and evaluation of a visual algorithmic music composition language for non-programmers

Matt Bellingham

Department of Music
University of Wolverhampton
matt.bellingham@wlv.ac.uk

Simon Holland

Music Computing Lab
The Open University
simon.holland@open.ac.uk

Paul Mulholland

Knowledge Media Institute
The Open University
p.mulholland@open.ac.uk

Abstract

Algorithmic music composition involves specifying music in such a way that it is non-deterministic on playback, leading to music which has the potential to be different each time it is played. Current systems for algorithmic music composition require the user to have considerable background knowledge of programming and/or music theory. However, much of the potential user population are self-taught music producers without the required background in either programming or music. To investigate how this gap between tools and potential users might be better bridged we designed Choosers, a prototype algorithmic programming system centred around a new abstraction (of the same name) designed to allow non-programmers access to algorithmic music composition methods. Choosers provides a graphical notation that allows structural elements of key importance in algorithmic composition (such as sequencing, choice, multi-choice, weighting, looping and nesting) to be foregrounded in the notation in a way that is accessible to non-programmers. In order to test design assumptions a Wizard of Oz study was conducted in which seven pairs of undergraduate music technology students used Choosers to carry out a range of rudimentary algorithmic composition tasks. Feedback was gathered using the Programming Walkthrough method. All users were familiar with Digital Audio Workstations, and as a result they came with some relevant understanding, but also with some expectations that were not appropriate for algorithmic music work. Users were able to successfully make use of the mechanisms for choice, multi-choice, looping, and weighting after a brief training period. The stop behaviour was not as easily understood and required additional input before users fully grasped it. Some users wanted an easier way to override algorithmic choices. These findings have been used to further refine the design of Choosers.

1. Introduction

Algorithmic composition typically involves structural elements such as indeterminism, parallelism, choice, multi-choice, recursion, weighting, and looping (Jacob, 1996). There are powerful existing tools, such as Max (Puckette, 1991) and SuperCollider (McCartney, 2002) for manipulating these and other elements of music. However, while these systems give great compositional power to musicians who are also skilled programmers (Wilson, Cottle and Collins, 2011), many musicians who are not also expert programmers find these tools inaccessible and difficult to understand and use (Bullock, Beattie and Turner, 2011).

This paper presents an evaluation of a prototype visual programming language (Bellingham, Holland and Mulholland, 2017) designed to allow structural elements of the kind involved in algorithmic music composition to be readily visualised and manipulated, while making little or no demand on programming ability. This system, called Choosers, centres around a novel non-standard programming abstraction (the Chooser) which controls indeterminism, parallelism, choice, multi-choice, recursion, weighting, and looping.

In this paper we present a programming walkthrough evaluation carried out with six pairs of undergraduate Music Technology students. The purpose of this evaluation is to:

- Test the ability of self-taught music producers without programming skills to use Choosers to carry out a range of rudimentary algorithmic composition tasks;
- Identify usability and user experience problems in the current design;
- Identify tensions and trade-offs in the interaction design of the system.

In the evaluation, pairs of participants were introduced to each element of the graphical programming language via short tutorial videos. Participants were given a range of practical tasks to complete on paper or a whiteboard. The evaluation facilitator played a Wizard of Oz role, rapidly translating participants' graphical solutions into runnable code that was fed into a non-graphical prototype version of Choosers so that participants could hear the musical results of their attempts.

2. Related work/problem setting

Various music programming languages are capable of algorithmic composition, although they require significant programming skills (Bullock, Beattie and Turner, 2011) and are therefore inaccessible to many users. Bellingham, Holland and Mulholland (2014) used the Cognitive Dimensions of Notations framework (Green and Petre, 1996) to review the usability of a representative selection of software capable of algorithmic music composition. The findings of the review included the following. First, we found that most existing software requires the user to have a considerable understanding of constructs in either graphical (e.g. Max, Pure Data) or text-oriented (e.g. SuperCollider, ChucK, Csound) programming languages: such knowledge requires a significant learning overhead. Second, users are often required to have an understanding of musical notation and/or music production equipment such as mixing desks and patchbays. Third, several programs imposed working practices uncondusive to compositional processes. Fourth, in some cases the user was unable to define, and subsequently change, the musical structure. Finally, complex visual design in graphical programming languages led to patches with multiple connections, making them difficult to read and to navigate.

3. Introduction to the system: Choosers

The following section provides a brief overview of Choosers, designed to cover enough detail to allow readers to understand the evaluation. The participants in the evaluation learned how the system worked by watching a sequence of short introductory videos and by investigating the system interactively via Wizard of Oz. For readers convenience the videos can be found online¹, and full details of the system design can be found in Bellingham, Holland and Mulholland (2017). The system has general musical expressivity, but for simplicity the present evaluation focuses on the manipulation of samples for algorithmic composition.

3.1. Samples and sequences

Samples are shown in boxes, and in the simulated interface can be auditioned by clicking on them. Samples can be assembled into sequences using arrows (see fig. 1). Samples in a sequence play in the order indicated by the direction of the arrows. Only a single arrow can enter or exit each element in a sequence. This deliberate limitation reflects the fact that parallelism and choice are dealt with elsewhere in the language. Boxes and sequences can be put inside other boxes, thereby packaging them into a single unit.



Figure 1 – Samples are shown in boxes, and a sequence is assembled via arrows.




Figure 2 – An annotated chooser.

4. Choosers: indeterminism, choice, parallelism and multi-select

Boxes referring to samples or sequences can be snapped together vertically to create what are known as Choosers (see fig. 2).

¹Available at <https://goo.gl/PFeAJf>

Fig. 2 shows a Chooser with two lanes, each containing a sample (drums and bass). The number 1 in the nose cone indicates that at run time, just one of the lanes will be selected at random (subject to restrictions described below). On different runs, different choices may be made. In fig. 2, by manipulating the number in the nose cone, any number of lanes from 0 to 2 can be chosen randomly to play simultaneously. A Chooser can have any number n of lanes. By manipulating the number in the nose cone, any number of lanes from 0 to n can be chosen randomly at run time and played simultaneously. Each lane has a weight associated with it. Consequently, in fig. 2, the drums are twice as likely to be chosen as the bass. Additionally, a weight of 'A' ('always play') can be used to ensure that the lane is always selected for playback.

4.1. Looping

Any sample can be set to loop indefinitely when selected on a particular run, or to play just once by the choice indicated in the status column (shown in fig. 2). Indefinite looping of a single sample may not always be desired, and for this reason we now introduce Time Choosers (see fig. 3).



Figure 3 – An annotated Time Chooser.

myVerse			
1	bass.wav	1	⌂
	drums.wav	2	↺
1	16 bars	1	×

Figure 4 – A full chooser.

4.2. Time Choosers and Full Choosers

If the Time Chooser shown in fig. 3 is attached onto the bottom of the Chooser shown in fig. 2, this produces the Full Chooser shown in fig. 4.

When the Full Chooser shown in fig. 4 is played, if the looping drums are chosen on a given run, they will not play indefinitely, but will be cut off after 16 bars by the Time Chooser. If the Time Chooser duration cleanly divides the sample duration, every repetition of the sample will run to completion, but if the Time Chooser duration does not exactly divide by the sample duration (for example, if the bass.wav sample in fig. 4 had a duration of 3 bars), the Time Chooser would cut playback mid-sample.

If, in a run of fig. 4, the non-looping bass were to be chosen, the bass sample would play once, and if this sample were less than 16 bars long, there would be silence until the end of the sixteen bars were reached.

4.3. Hard vs. soft stop of loops

If the status column in the time chooser is set to $>$ (indicating a soft stop) rather than \times (indicating a hard stop) then, when the time chooser ends, the sample will continue to play until the end of its current loop.

4.4. Terminology: Soundable Choosers, Time Choosers and Full Choosers

Now that Time Choosers and Full Choosers have been introduced, in order to avoid ambiguity, we will refer to Choosers with no attached Time Choosers, such as those shown in fig. 2, as Soundable Choosers.

4.5. Time Choosers on their own

A Time Chooser can be used alone as part of a sequence – however, when used in this way it will simply result in a rest of the specified duration.

4.6. Time Choosers with multiple lanes

More generally, the purpose of a Time Chooser within a Full Chooser is to moderate in a non-deterministic manner how long the Soundable Chooser and its individual lanes play. Possible interactions between the settings of soundable and Time Choosers can make the results more varied than might be imagined.

A Time Chooser's nose cone can be set to either one or zero. If set to one, one time lane will be chosen at run time. If it is set to zero no time lanes will be selected and the Soundable Chooser will run as though there is no Time Chooser. This allows for quick low viscosity arrangement changes, with the possibility of infinite playback if the Soundable Chooser lanes are set to loop. If the Soundable Chooser is not set to loop, the sample(s) will play and the Chooser will be released when they have finished playing, regardless of length.

5. Method

5.1. Participants

Seven pairs of undergraduate Music Technology students took part in user tests utilising a Wizard of Oz prototyping methodology. These users were targeted as they are typically neither programmers nor traditional musicians. While they may be conversant with some elements of music theory, the predominant background is self-taught music producers with experience of making music electronically using music sequencers/DAWs. The users were introduced to each element of the graphical programming language via short tutorial videos². Users were given a range of practical tasks to complete on paper (see fig. 5) or on a whiteboard, and their outputs were played by the facilitator using a set of SuperCollider (McCartney, 2002) classes written to implement the musical abstractions behind the system. The user tests were videoed and transcribed to assist in the analysis presented here.

All participants were asked to complete a short questionnaire before taking part in the user tests. Of the fourteen participants all were musicians, and ten had some formal training. Six participants did not read any music notation: of those that did, most could read common music notation as well as chord notation. All participants were familiar with DAWs, with Logic Pro (Apple Inc., 2013) mentioned by all fourteen users. Other DAWs mentioned included Pro Tools (6 mentions), Cubase (2 mentions), FL Studio (5 mentions), Reason (1 mention), and Ableton Live (1 mention). Pure Data (Puckette, 1997) a visual audio programming language, was mentioned by two participants. Twelve participants had experience using hardware for music performance HCI tasks (such as drum pads or control surfaces). The participants were not habitual performers; seven of the fourteen participants do not perform with or for others. Of those that do, three perform in church, and four occasionally play with friends in private. Five of the fourteen participants felt they had some experience in computer programming. Of these, two considered writing for the web (HTML and CSS) to be programming. Given that HTML/CSS development is focussed on markup and layout we can see that more than two thirds of participants do not have experience of writing algorithms in a programming language. One participant listed the use of Pure Data and SuperCollider. Eight of the fourteen participants did not know what algorithmic music was at the start of the user tests. Four participants felt that they knew what algorithmic music was, but had not created any.



Figure 5 – Participant Chooser work using paper.

1	drums.wav	1	
	bass.wav	1	

Figure 6 – Scenario one.

5.2. Walkthrough protocol

Participants were asked to take part in eight scenarios, as reproduced below. The users were free to discuss the work and to ask for clarification with the administrator of the test. Users were asked to act as active participants in the research, and to help in categorising any issues that were raised The

²Available at <https://goo.gl/PFeAJf>

categorisations that users were asked to use –taken from the programming walkthrough method (Bell, Rieman and Lewis, 1991; Bell et al., 1992) – were questions (e.g. why does the loop do that?), problems (e.g. I don’t understand what these lanes are for), suggestions (e.g. maybe the cone should be a different shape), and other observations (e.g. I like the fins). In addition, participants were asked if they could think of any other ways in which each scenario could be completed. This prompted a discussion on alternative routes in order to test understanding and to capture user expectations.

5.3. Walkthrough scenarios

We now present the eight scenarios issued as part of the user tests. The videos, images, and verbal instructions issued to participants are shown in boxes. This section includes a brief overview of the results from each scenario. A more detailed reflection on the design issues is in sec. 6.

5.3.1. Scenario 1 - understanding the Soundable Chooser

Participants were shown a short video³ which introduced lanes, nose cone, weighting inc. always play, and loop/non-loop functionality.

Here is a Soundable Chooser with two samples.

- If this Chooser is played by itself, how many samples will play?
- How do you know?
- How likely is it that the drums.wav sample will play? How could you make it more likely to play?
- How could you make the Chooser play both samples?
- How would you make it play no samples?

This scenario prompted a number of clarifying questions from participants, all of which could be answered simply by playing the video again. All participants were able to complete this scenario without error.

5.3.2. Scenario 2 - creating a Soundable Chooser

Make a Soundable Chooser which has three lanes – those lanes should contain looping drums, bass and guitar samples. Make it so that two play at once – the drums always play, and either bass or guitar will be selected with equal probability.




Figure 7 – Paper template given to participants for scenario 2.

1	2 bars	1	X
	4 bars	2	X

Figure 8 – Scenario 3.

- Make it so the guitar doesn’t play

The Soundable Chooser design was understood quickly by most participants. The layout of the Chooser – vertically stacked lanes – prompted some interesting assumptions from participants in two groups. The participants appeared to apply existing knowledge of horizontally-scrolling DAWs, such as Pro Tools or Logic Pro, in which vertically stacked lanes play concurrently. However, in the Chooser design, such lanes contain material which may or may not be selected for playback and concurrency cannot be assumed without also considering the number in the nose cone, the lane weight, and the “always” feature described below. This assumption, made twice in the user tests, was corrected at the start of the test and users did not encounter any further problems.

³Available at <https://youtu.be/gFbWz-F-WmE>

The use of 'A' ('always play') in the weight column of Soundable Chooser lanes was understood quickly by all participants. Users were asked to prevent the guitar from being available for selection, and all groups chose to change the guitar's weight to zero. This answer was correct, but was not the only way to achieve the desired result. One participant correctly suggested that the bass could also be set to A, meaning that the guitar would not be selected given the restriction imposed by the number in the nose cone.

5.3.3. Scenario 3 - understanding the Time Chooser

Participants were shown a short video⁴ which introduced Time Choosers, multiple lanes, and nose cone restrictions.

Here is a Time Chooser with two lanes.

- Describe what will happen when this Chooser is run by itself.
- The nosecone is currently set to 1. What else could it be set to? What would happen if it is changed?
- How could you make a duration of 2 bars most likely to be selected?

Two groups found multiple vertically aligned lanes confusing, as in the previous scenario. One participant assumed that both durations would play together, presumably due to the behaviour learned through the use of DAWs. Another user thought that the first duration would play, followed by the second duration. In both cases the participants were shown the relevant section of the tutorial video again, and they then fully understood the Time Chooser mechanism.

5.3.4. Scenario 4 - creating a single-lane Time Chooser

Keeping the Soundable Chooser you made for scenario 2, create a single-lane Time Chooser.



Figure 9 – Paper template used for scenario 4.

2	drums.wav	A	↻
	bass.wav	1	↻
	guitar.wav	1	↻
1	4 bars	1	X

Figure 10 – The result of the final task in scenario 4.

- Make a four bar rest.
- Now you have made the rest, find two ways to quickly skip it. What do you think the nose cone value could be? What would happen if you gave the nose cone that value?
- What impact would there be if you changed the weight of the lane?

Participants were shown a video⁵ which introduced Full Choosers, duration control via the Time Chooser, and the hard and soft stop mechanism.

- Next, take the Time Chooser and snap it onto the Soundable Chooser created in scenario 2. What is the impact of this?

The user tests were designed to introduce a Time Chooser as playing silence for a given duration or, from an alternative viewpoint, introducing a musical rest. We subsequently moved on to introduce the

⁴Available at https://youtu.be/p03--FbA_r0

⁵Available at <https://youtu.be/8AWcol0MMwY>

Time Chooser’s fundamental role, which is to provide a mechanism for constraining the duration of a Soundable Chooser. We found that introducing Time Choosers in this way created some confusion among users, and future user testing will reverse the order in which the two functions are introduced. Four of the six groups found that the rest functionality of the Time Chooser was not obvious, although all participants were able to use it effectively after discussion. When the user test moved on to the use of a Time Chooser to control a Full Chooser’s duration one participant interpreted it as a rest – they continued to use the logical framework of rests rather than re-contextualise the functionality in the Full Chooser. Tellingly, a significant number of participants guessed that Time Choosers would be used to control the duration of Soundable Choosers before the functionality was introduced. We consider this in more detail in sec. 6.2.

Conceptually, hard and soft stops were understood immediately by half of the participants. These users were able to make musical use of both stop types. Those who did not initially understand the difference were almost all able to clarify the stop behaviour via a short discussion. One participant was surprised by hearing the effects of a soft stop in his group’s final scenario work (‘why isn’t it stopping?’), prompting another conversation. The stop functionality is further explored in sec. 6.1.

Two groups felt that, while the icon for a hard stop was understandable, the icon for a soft stop did not make its function clear. When the intention behind the icon choice was explained (× as a hard stop as it mirrors a ‘stop’ road sign, > as a soft stop as it looks like an arrow to allow the material to continue) some participants felt that it was more understandable. One participant remained unsatisfied with the soft stop icon, although they could not offer an alternative.

Just as the icon choice had an effect on learnability, the language used for the stops also seems to have been significant. For those participants who did not immediately understand the stop functionality we explained hard stops as ‘rude’ and soft stops as ‘polite’. These were the original names used in early development, and they were useful in allowing the participants to contextualise the stop behaviour. It is not clear from this process whether one name is more descriptive or learnable than the other, as not all participants were introduced to these terms and those who were, encountered them as a secondary adjective term to support the ‘hard’ and ‘soft’ terminology.

5.3.5. Scenario 5 - understanding a Full Chooser

Look at this Chooser and say what will happen when the Chooser is played. The drums and bass samples are four bars long. The marimba sample is two bars long.

drums.wav	3	↺
bass.wav	2	↺
marimba.wav	1	↺
8 bars	1	X

Figure 11 – Scenario 5.

guitar.wav	A	↺
BVs.wav	A	↺
vocals.wav	A	↺
8 bars	1	X
16 bars	2	>

Figure 12 – Scenario 6.

- How many samples will play? Will they loop or play once? What effect would changing the loop setting on the drums have?
- How long will the Chooser play for? What happens when the duration elapses?
- What would happen if the Time Chooser was set to a soft stop?
- What would happen if the Time Chooser nose cone stayed at 1 and the Soundable Chooser nose cone was changed to 2? To 1? To zero?
- How could you make it infinite playback? How could it be made into a rest? Skipped entirely?

When asked to produce a Full Chooser which allows for infinite playback, the users were expected to set the Time Chooser’s nose cone to 0 in order to prevent it from controlling the Soundable Chooser’s

duration. One user wanted to be able to set the time lane's duration to ∞ . While not one of our options, this would potentially be a valid input. Two participants wanted to use the 'A' (always play) mechanism to set infinite playback – their logic was that they wanted to override the set duration. Another user wanted to be able to loop the Time Chooser to create infinite playback.

One user wanted Soundable Choosers to segment to show the number of loops in a given duration. For example, if a duration of 8 bars was selected, and the upper soundable lane contained a 2-bar sample and the lower lane contained a 4-bar sample, the upper lane would show 4 segments and the lower lane would show 2 segments.

One group were interested in visual feedback from the Choosers. They wanted to have a progress bar, which presents some interesting trade-offs and challenges to the system. This is explored in sec. 6.1 and sec. 6.2.

5.3.6. Scenario 6 - understanding a Full Chooser with multiple lane times

Here is a Chooser containing a Time Chooser with multiple lanes.

- What do you expect to happen in the Soundable Chooser?
- What will happen in the Time Chooser? Which lane is more likely to be selected? What are the consequences of the selection of the uppermost Time Chooser lane? What will be different if the lower Time Chooser lane is selected?
- What other values are possible for the nose cone of the Soundable Chooser?
- What other values are possible for the nose cone of the Time Chooser?

This scenario, which introduced multiple Time Chooser lanes, was incorrectly explained by two user groups (see sec. 5.3.3) who did not understand the lane selection mechanism. This error may be linked to the application of an understanding of horizontally-scrolling DAWs (see sec. 6.3). The use of numbers for multiple parameters may also have made the functionality unclear (see sec. 6.5).

This scenario asked users to consider the possible nose cone values given three Soundable Chooser lanes with a weight of 'A' (always play). The correct answer (either three to play all lanes, or zero to skip the Soundable Chooser) was quickly identified by most participants. Two groups questioned the semantic logic of allowing the 'always play' directive to be overridden and, if it can be overridden, why this can be only performed in one context. This was an interesting example of users applying their knowledge from within the system and is considered in more detail in sec. 6.3. The issue is indeed problematic and led to the infinity option discussed in more detail in sec. 7. The infinity option proposed in sec. 7 does not require the nose cone to be changed or constrained, while allowing similar semantic meaning.

5.3.7. Scenario 7 - creating a Full Chooser

Using the templates (provided on paper), create a Full Chooser which:

- Has four soundable lanes, three of which will play at any given time. Drums and bass, which always play, and are set to loop: guitar and vocals, where the guitar is twice as likely as vocals to be selected for playback. Neither should loop.
- Has three possible durations, of which one will be selected – 2 bars with a hard stop, 4 bars with a soft stop, and 5 bars with a hard stop. Make the 2 bar duration twice as likely to be selected as the 4 and 5 bar durations.

This task was performed quickly and accurately by all groups. The users in two groups requested preset Choosers to enable them to build a viable piece of music quickly. Presets could also provide tutorial support by providing a framework around which users could experiment.

5.3.8. Scenario 8 – playground

Users were shown the final video⁶ which introduced the sequence mechanism.

⁶Available at https://youtu.be/bBnngW-W_HU

Using the templates and samples available, make a piece of music which uses a sequence of three Choosers. The music will be recorded and shared online. The piece should be musically satisfying even if it is run only once. If it is run more than once it should be different in some way.

The participants were able to create some interesting musical material, and all enjoyed hearing the results of their work. This was the first time that any of them had created algorithmic music. Three participants felt that the user interface used numbers for too many parameters. Upon questioning, their concern was that there was insufficient differentiation between quite different user interface elements, making the interface both difficult to learn and confusing in operation. Numbers are used in the nose cone, the time lanes (e.g. '8 bars'), and the weight column. We consider this in sec. 6.5.

One suggestion was to change the Time Chooser nose cone to an on/off icon given that the only legal options were zero and one. It is interesting to note that some users were able to identify that a Boolean value is required here, and we consider the implications of this in sec. 7.

Most users used this final task to test their understanding of the system, rather than creating a cohesive piece of music. Of the seven groups, one decided to create a piece which developed thematically. To do this they wanted to reuse a Chooser, making minor changes to it. We discussed copying the original Chooser, and this process met their requirements.

One group was interested in reusing some elements of the composition created in the final scenario, and so we discussed the nesting functionality which was left out of this round of user testing (see Bellingham, Holland and Mulholland (2017) for an explanation of nesting in Choosers). They understood the concept and it entirely met their requirements.

5.3.9. Final questions

At the end of the user test, participants were asked the following three questions:

- Can you see anything this would be useful for?
- Can you see any ways in which this is similar to other tools you have used?
- Is there anything that is made easier by this system? Anything which was not possible made possible/hard and made easier?

Three groups commented on the 'boring' design. The layout of Choosers was not seen as problematic, but some users wished for a more stylish and polished presentation. Two groups requested colours to enhance usability: within one group, one user wanted automatic colour selection (denoting lane type) and the other user felt that user-controlled colour selection would better support sorting and arrangement tasks. Multiple users were interested to know if lanes could be rearranged to visually organise lanes into instrument groups. One user suggested that lane arrangement could be an alternative to the weight column — moving a lane higher would result in a higher probability of playback. This is similar to one mechanism which was considered and rejected before the user tests: it was replaced by the weight column as the column allows for multiple identical lane weights, quick auditioning, and user-controlled lane ordering to assist with musical arrangement. One user requested instrument icons for soundable lanes, partly in response to being unaware of the marimba (one of the samples used in the user test).

6. Reflection on design issues

The findings from the user tests outlined above have various implications for the design of Choosers.

6.1. Musical issues

Repeating phrases, and the musical interaction between phrases, are crucially important in a music system. These have therefore been brought to the surface via the loop and hard/soft stop behaviours. We found that the stop behaviour was confusing to four of the seven pairs of users, and the documentation will be enhanced to better explain the system. The hard and soft stop system can be thought of in a number of ways. For musicians, one useful way is to consider soft stops as suitable for melodies, and hard stops for accompaniment. Melodies are therefore allowed to finish, whereas accompanying

elements are stopped when the duration of the Chooser elapses.

Three of the fourteen users were keen to have a visual indication of current position with respect to duration, such as a progress bar. While this seems a reasonable request, it is complicated by the non-deterministic nature of the system.

A significant number of participants found the use of Time Choosers for both rests and Chooser duration to be confusing. This was largely due to rests being introduced before the Time Chooser's primary function, which is to control the duration of a Chooser.

None of the participants had experience in algorithmic composition, so these sessions essentially introduced algorithmic compositional tools while testing the interface. This led some participants to presume that the concepts themselves were novel. Some time was spent discussing the desirability of algorithmic processes rather than this specific implementation. Two participants assumed that the process would lead to a linear audio file, which indeed it can, but many use cases would require the music to remain nonlinear. Future evaluations could explore the nature of the resistance to nonlinear playback, including how it is related to expectations set by commercial music creation software and linear playback. We are also interested in the use of Choosers in genres which routinely incorporate extemporaneous changes and improvisation, such as folk and jazz.

One group specifically wanted a mechanism to allow them to easily reuse material for thematic development. The design of Choosers allows for this via nesting, although it was not included in the user tests for simplicity. The users were shown nesting in response to their questions and found it to meet the need they had expressed.

Choosers can be used in the creation of a range of music. However, given the unusual combination of usability and affordances, Choosers are particularly suited to music in which users would benefit from easy access to non-linear playback. Some classic Minimalism techniques (Potter, 2002), such as phasing (Scherzinger, 2005), are easily achievable using Choosers. Game music is often non-linear, created using layers of musical material which are triggered by in-game events (Collins, 2008). Such material can be created using Choosers, and we have a mechanism which would allow for external input via OSC or an alternative protocol; this would allow a game engine to trigger changes in the music. Choosers also allow musicians and music producers to create nonlinear versions of existing recordings by loading alternate takes into Choosers. The playback could range from very close to the original (e.g. algorithmically switching between vocal takes of the same melody) to playing significantly different material (e.g. branching to play different sections), depending on the decisions made by the creators.

6.2. Programming-related issues

As shown in sec. 3, the Soundable Chooser nose cone slopes down and the Time Chooser nose cone slopes up – this allows them to be joined together and communicates the required upper/lower order to the user. Interestingly, some users guessed the combination of Soundable and Time Choosers, suggesting that the nose cone shapes of the two Chooser types were effective in communicating their combinatorial usage.

The Chooser system is designed to allow for consistent logic to be applied across Soundable and Time Choosers where possible. Participants in the user tests successfully reused elements of the Soundable Chooser system when manipulating duration, but there were some cases where such reuse or re-contextualisation was not possible. Interestingly, the actions of the users in these cases would have made sense neither from a musical nor programming perspective, but the rationale behind these requests is instructive as it shows how users understand the tools in the system. For example, in scenario five (sec. 5.3.5) two participants wanted to use the 'A' (always play) mechanism to set infinite playback – they wanted to override the set duration and had understood 'A' to be a global override control. In a similar example, one user wanted to be able to loop a Time Chooser. If the system were to be changed to allow for a set number of repeats, rather than an infinite loop, such a move may be desirable.

Users will also require access to metadata – for example, to check the length of a sample loaded into

a soundable lane in a Chooser. Such metadata could be shown via a tooltip, accessed by hovering the mouse over a lane.

6.3. Shared and existing knowledge

One design motivation is to enable people to understand the system very quickly. The Chooser design tacitly draws on a number of systems of existing knowledge.

Some users wanted to be able to leverage their existing understanding of DAW software and found it frustrating that they needed to learn new paradigms for duration, synchronicity, and so on. This is an example of technological framing (Orlikowski and Gash, 1994). The knowledge gained by using other music software can be useful, but it can also prove problematic if the design of the software being learned is sufficiently differentiated. As a result, there is much to be gained by following standard design conventions where possible, as this maximises the user's ability to reuse existing knowledge. One interesting example was seen in scenario 6 (sec. 5.3.6), in which one pair of users learned the rules of Choosers and then wanted to use the same rules elsewhere.

Technological framing, and the expectations set by the use of commercial DAWs, may be an influence on user requests for a progress bar and the conversations on the desirability of nonlinear vs. linear playback that took place during the user tests (as considered in sec. 6.1).

6.4. Metaphor

Interface metaphors are very common and can be useful in communicating the roles of the software and setting realistic expectations when users are familiar with the original interface. However, such metaphors can become problematic if users are not familiar with the original interface.

Related to technological framing is the assumption, ubiquitous in Digital Audio Workstations, that signal flow and processing will be applied using a mixing desk metaphor. Such virtual desks often make use of skeuomorphism (such as the fader caps and rotary potentiometers in Pro Tools), although some other designs have made graphical changes while retaining the overall layout. As an example, Ardour's use of a textured 'strip' instead of a fader is still skeuomorphic as it makes use of a ribbon controller metaphor but, in an attempt to improve mouse control by increasing the size of the target, it does not follow the traditional desk layout.

Given that DAWs are now capable of performing all mixdown tasks, and the financial cost of consoles and outboard effects processors can be prohibitive, many users learn in a virtual studio environment rather than on hardware. Many DAWs were designed to mimic hardware in order to leverage existing knowledge and ease the transition from hardware to software. However, now most people are introduced to music production via software, and many do not use hardware, there is an opportunity to revisit some design assumptions.

Some users felt that the user interface was 'boring', lacking the use of colour, metaphor, and skeuomorphic design common in DAWs. This may be another example of technological framing (Orlikowski and Gash, 1994). We can also consider the impact of metaphor in music software by making use of the Cognitive Dimensions of Notations (Green and Petre, 1996). Using this framework, the closeness of mapping and role expressivity of a mixing desk can be implied by making a software recreation look and function like hardware.

Some users had difficulty understanding the outcome of hard and soft stops in Choosers. The vast majority of music production software is focussed on the creation of linear music, and the concept of 'play until finished' is rarely implemented. As a result, none of the user test participants had encountered it, and did not have a frame of reference for why it might be desirable. As a result, there is not a clear existing metaphor for what we refer to here as a 'soft' stop. Users agreed that the × icon represented a traffic stop sign and that it was a suitable analogy for 'stop now', but the > icon used for a soft stop was not immediately understood as there is no readily accessible metaphor.

6.5. Arithmetic

The use of numbers and arithmetic relationships in an interface can be a valuable organising tool, as they are more or less universally familiar and can concisely represent many relationships. The decision to use numbers for several parameters was motivated by parsimony and consistency. However, the use of numbers for multiple parameters was perceived as negative by three participants. Upon questioning, the issue was that numbers meant different things in different parts of the interface. The Chooser design presented to participants in the user tests made use of integers in five different ways: for the number of simultaneously playing soundable elements, weight, duration, repeats, and Time Chooser on/off. Despite this, for different reasons, the user issues surrounding the ‘always play’ option led us to consider extending the range of numerical concepts used in the interface, by allowing the metaphorical use of ∞ as a weight (to outrank any positive integer weight) as discussed in the next section.

In sec. 7, we propose changes to Chooser design to address these various issues.

7. Design problems and candidate solutions in Choosers

Given the problems for some users with the use of integers for multiple parameters (see sec. 6.5), we propose the use of a simple on/off icon for the Time Chooser nose cone. Interestingly, one pair of users suggested this change in the user tests. Scenario five, outlined in sec. 5.3.5, showed that two users wanted to leverage the ‘always play’ mechanism beyond the weight column, and one user wanted to set the duration of a time lane to infinity. We propose a change to Choosers which allows for both mechanisms.

We propose a design change which allows the user to allocate a maximum possible weight (∞) for a lane, thereby guaranteeing that it will play if the nose cone number is high enough to allow all such lanes to play. When allowing ∞ as a weight, a useful metaphor is to think of lanes with weight ∞ as having paid for “priority boarding”, as when boarding an aircraft. Lanes with weight ∞ will always be chosen before any lanes with any finite weight. Compared with the ‘always play’ mechanism, this has the potential for greater clarity when the number of maximally-weighted lanes exceeds the number of the nose cone. In such cases, under the current ‘always play’ system it is not obvious whether ‘always play’ should override the nose cone or vice versa. Under the proposed system, the nose cone would determine the number of lanes to play, and if that was less than the number of lanes with weight ∞ , the winners would be chosen from those lanes at random. We are also considering the use of a maximum value (∞) for the nose cone of a Soundable Chooser (‘play all available lanes’) and for the duration of a time lane (‘play forever’).

We propose that future work will introduce Time Choosers in the context of a Full Chooser, with the rest functionality introduced later as a special case. Tutorial materials will provide a clear explanation and will offer context and examples. The value of all of these proposed changes will be tested empirically.

8. Conclusions

Choosers were developed to allow non-programmers access to algorithmic composition tools and processes. The design principles were to leverage parsimony to enhance learnability; to surface musically meaningful actions, and to make them quick and easy; to allow both bottom-up and top-down construction; and to make use of progressive disclosure to allow for advanced use without harming usability for beginners.

The user tests outlined here show that non-programmers were able to successfully use Choosers to create a number of short pieces of music. Future work will focus on the refinement and re-evaluation of the Chooser notation and supporting materials.